

Gibbs Sampling for the Multivariate Gaussian mixture model

Computational Statistics Project 3

Leonardo Ruggieri

1 Introduction

Mixture models are simple yet powerful instruments when it comes to capture complex distribution and model a wide variety of data. One common task is to perform clustering by means of a hierarchical model with latent variables. The basic element is the mixture of Gaussian distributions, that considers a superposition of Gaussian densities according to some prior mixing coefficients. In the following, a simple Gibbs sampler is motivated and implemented with synthetic data in a Python code, for which convergence diagnostics is provided. Finally, an application to color image segmentation is delivered.

2 The model

2.1 The Gaussian mixture distribution

The Gaussian mixture model provides a richer class of density models than a single Gaussian distribution. The mixture distribution can be defined as follows:

$$p(x) = \sum_{z=1}^K \pi_z \mathcal{N}(x|\mu_z, \Sigma_z)$$

where π_k are the mixing coefficients [1]. Given that their sum must be equal to one, that is, $\sum_{z=1}^K \pi_z = 1$, and that $\mathcal{N}(x|\mu_z, \Sigma_k)$ is non-negative, $0 \leq \pi_k \leq 1 \ \forall k = 1, \dots, K$ must hold, so that we have $p(x) \geq 0$ on its support. Hence, the mixing coefficients satisfy the requirements to be probabilities.

2.2 Formulation in terms of latent variables z

We can equivalently represent the mixture distribution by means of a latent (unobserved) variable z . Hence, with multiple data points x_i , there is a corresponding latent variable z_i for each one of them, which identifies its cluster assignment. This latent variables will be critical in our applications: we want to make Bayesian inference on it, so to discover which group every data point belongs to. Thus, the marginal density will be:

$$p(x) = \sum_z p(z)p(x|z)$$

2.3 The model

By specifying the priors for the cluster means and by fixing a common variance σ^2 for the observations, the hierarchical model is the following:

$$\mu_z \sim \mathcal{N}(0, \lambda^2 I_d) \quad z = 1, \dots, K \quad (1)$$

$$z_i \sim \text{Categorical}(1/K, \dots, 1/K) \quad i = 1, \dots, n \quad (2)$$

$$x_i | z_i, \mu \sim \mathcal{N}(\mu_{z_i}, \sigma^2 I_d) \quad i = 1, \dots, n \quad (3)$$

The mean parameters are independently drawn from a common prior, expressed by a multivariate normal centered in zero. In addition, d is the dimension of the space of the observation x_i 's. Moreover, μ_{z_i} is the mean vector of the cluster to which x_i belongs.

3 Gibbs sampling

3.1 Motivation

The MCMC approach, in general, requires to simulate a Markov chain that is invariant with respect to the target distribution. The Gibbs sampler is a technique for indirectly generating random variables from a marginal distribution without having to calculate the density explicitly[3]. Specifically, it can be useful when dealing with multidimensional target distributions, for which the analytical formulation can be hard to compute. With the Gibbs sampler, we only deal with full (or complete) conditional distributions, for which it is typically easier to obtain. Then, the Ergodic theorem, the very same realizations of the Markov chains are used to obtain an approximation of the posterior distribution and the Bayesian estimates of interest. The Markov chain is built by using the complete conditionals as transition densities: thus, each of the generated subsequences is a Markov chain, which has the corresponding marginal as invariant distribution. In general, the process is as follows:

1. Initialize $X_0 = x_0$
2. For $t=1, 2, \dots$, given $\mathbf{x}=(x_1^{(t)}, \dots, x_d^{(t)})$ draw:

$$X_1^{t+1} \sim f_1(x_1 | x_2^{(t)}, \dots, x_d^{(t)})$$

$$X_2^{t+1} \sim f_2(x_2 | x_1^{(t+1)}, x_3^{(t)}, \dots, x_d^{(t)})$$

...

$$X_d^{t+1} \sim f_d(x_d | x_1^{(t+1)}, \dots, x_{d-1}^{(t+1)})$$

Hybrid Gibbs sampling The Gibbs sampler is a Metropolis-Hastings algorithm, where the acceptance probability of each step is equal to 1. The two methods are often used combined: for instance, when the full conditionals of all the variables and parameters of interest are not analytically tractable, the

Metropolis-Hastings can be incorporated within a Gibbs sampler to draw samples from the variables whose full conditionals cannot be analytically determined. This scheme is often called *Metropolis-Hastings-within-Gibbs* algorithm. One example of the use of such algorithm is the *latent position cluster model* for social network [4], which builds on our hierarchical model: the observations are mapped into a latent space, whose coordinates are drawn from a Gaussian mixture.

3.2 Gibbs sampling for the Gaussian mixture model

Let's go back to our original mixture of multivariate Gaussians. For a sample of size n , the joint density of x , z and μ , where z and μ are the latent variables, is as follows:

$$p(x, z, \mu) = p(\mu) \prod_{i=1}^N p(z_i) p(x_i | z_i, \mu)$$

In order to get the posterior distributions of μ , we wish to compute $p(\mu | x) = \frac{p(\mu, x)}{p(x)}$. In order to do so, we need to work out the marginal distribution $p(x)$. However, such a distribution, shown in 4, is rather intractable from a computational standpoint, as one must consider all the possible configurations of cluster assignment, with the computational complexity increasing exponentially in K [2].

$$p(x) = \sum_z \int p(\mu) \prod_{i=1}^n p(x_i | z_i, \mu) d\mu \quad (4)$$

We set a uniform prior over the cluster assignment for each data point, so that $\pi_{z_i} = \frac{1}{K}$. In order to obtain the Gibbs sampler, the full conditionals must be worked out. As far as the complete conditionals of z are concerned, by exploiting the conditional independence, one can easily get:

$$p(z_i | \mu, z_{-i}, x) = p(z_i | \mu, x_i) \propto p(z_i) p(x_i | \mu) = \pi_{z_i} \phi(x_i; \mu_{z_i}, \sigma^2 I_d)$$

where ϕ is a Gaussian density evaluated at x_i and μ_{z_i} is the mean vector. Similarly, one can obtain the complete conditional for μ_k :

$$p(\mu_k | \mu_{-k}, z, x) = p(\mu_k | z, x) \sim \mathcal{N}(\hat{\mu}_k, \hat{\lambda}_k)$$

where $\hat{\mu}_k$ and $\hat{\lambda}_k$ are easily obtained by conjugacy results.

The final algorithm is shown in the pseudocode below.

Algorithm 1: Gibbs Sampling algorithm for Gaussian mixture

Initialize mixture locations μ and mixture assignments z .

```

for  $t = 1, \dots, T$  do
  for  $i = 1, \dots, n$  do
    for  $z = 1, \dots, K$  do
      | Compute  $c(z_i = z) = \pi_z \phi(x_i; \mu_z, \sigma^2 I_d)$ 
    end
    for  $z = 1, \dots, K$  do
      | Normalize  $p(z_i = z) = \frac{c(z_i = z)}{\sum_{z=1}^K \pi_z \phi(x_i; \mu_z, \sigma^2 I_d)}$ 
    end
    Draw  $z_i | \mu, x_i$  from a Categorical with weights  $p(z_i)$ 
  end
  for  $k = 1, \dots, K$  do
    Compute  $n_k = \sum_{i=1}^n z_i^k$ . Compute  $\bar{x}_k = \frac{\sum_{i=1}^n z_i^k x_i}{n_k}$ 
    Compute  $\bar{\mu}_k = \frac{n_k / \sigma^2}{n_k / \sigma^2 + 1 / \lambda} \bar{x}_k$ . Compute  $\hat{\lambda} = (n_k / \sigma^2 + 1 / \lambda^2)^{-1}$ 
    Draw  $\mu_k | z, x$  from  $\mathcal{N}(\bar{\mu}_k, \hat{\lambda}_k)$ 
  end
end

```

4 Application on synthetic data

4.1 Description

We now turn to fit the model to some generated data. In particular, we generate 50 observations coming from each of the 3 different bivariate Normal distributions, respectively centered in the origin, in $(1, 2)$ and in $(3, -4)$ of the Euclidean space, with the same isotropic covariance matrix, as shown in 1.




Fig. 1: Synthetic data coming from 3 different clusters.

4.2 Ergodicity and convergence diagnostics

Thanks to the convergence results, we know that, eventually, sampling a value from the generated Markov chain is equivalent to sampling a value from the invariant distribution. In addition, thanks to ergodicity, one can take more values from the same Markov chain as to obtain a sample from the distribution of interest.

However, one may want to check if the hypotheses of such results are true. In particular, the convergence to the stationary distribution, the convergence of averages and the convergence to i.i.d. sampling must be assessed. In the following, trace plots and tests are provided after 1000 iterations of the algorithm.

Trace plots In Figure 2, the trace plots for the two component of the mean parameter for one of the cluster is shown. By graphical inspection, the chains seem to have reached stationarity. However, more formal checks are provided below.




Fig. 2: Trace plots of the means of one cluster.

Effective sample size On the basis of the autocorrelation function, a Python function that computes the effective sample size has been implemented. The effective sample size is the number of independent Monte Carlo elements to get the same precision of the MCMC procedure [5], and it is computed according to Equation 5. In particular, the effective sample size oscillates between 320 and 340, suggesting that the number of iteration run might be sufficient for our purpose.

$$s_{eff} = \frac{s}{1 + 2 \sum_{l=1}^{\infty} \text{corr}(X^t, X^{t+l})} \quad (5)$$

Kolmogorov-Smirnov test The Kolmogorov-Smirnov test is a nonparametric test for stationarity: it takes two samples and checks whether the two are generated from the same distribution. In our case, after a *thinning* procedure to reduce the autocorrelation of the chains [6], we divide the chains in three parts: the first part can be considered the *burn-in* phase; the second and the third segments are given as input to the test. This tests provides an evidence supporting or rejecting the stationarity of our chains.

The test has been conducted on the chains of μ and z at the end of the procedure. The p-value of the tests conducted, whose null hypothesis is that both chains come from the same distributions, are always high and above 30%, providing another evidence that the chain has reached stationarity. Similar results show for the cluster assignments.

4.3 Results

We consider the first third of the chain as the *burn-in*. Hence, we take the successive values in order to obtain the Bayesian estimates for the cluster assignments and for the mean vectors. In particular, we take the average of the values obtained from the chain and we use these values as our estimates, relying on the Ergodic theorem.

In the following graph, after having obtained the estimates of the variables of interest, 50 points for each cluster were sampled according to the generative process.




Fig. 3: Data generated with the estimated parameters (1000 iterations).

4.4 Implementation issues

Some issues may arise on the choice of the starting values of the chains. Indeed, this aspect has an impact on the mixing of chain, ultimately. Sometimes one may choose to start from the Maximum Likelihood estimate. another choice is to start from the prior: the cluster assignments have been initialized with values drawn from the prior distribution. Nevertheless, we can notice that even with different starting values, we obtain very similar results at the end of the procedure.

Another issue may arise in the evaluation of the convergence: the methods implemented above indicate quite clearly that the convergence has been reached. However, these kind of results are not always straightforward to interpret. One can then decide to deploy further tools based on the comparison of multiple chains, such as the Gelman-Rubin convergence diagnostics. However, such methods might be computationally demanding, as multiple chains need to be generated for each parameter. For more complex application, like the following, the opportunity cost of running such a test is high: indeed, one might simply prefer to run the chain for a longer time.

The Gibbs sampling is a very convenient method to deal with high-dimensional densities. However, particular attention should also be provided to the hyperparameters specification, which may severely affect the mixing time of the chain and its convergence.

5 Application to Image Segmentation

The second application is to show the great potential of this simple model. Image segmentation algorithms aim at partitioning a given image into homogeneous pattern classes. In the context of our model, this task is accomplished by assigning each pixel to one of the K possible classes, which corresponds to one of the K different mixture components. [7] [8]

5.1 Description

An image of 32×32 pixels has been provided as our data set. Each pixel is characterized by 3 dimensions, that express the intensity of each of the RGB channels. For each channel, there are 256 possible color intensities, for a total of 256^3 possible combinations. In our application, 10 clusters are specified, which means that all the 256^3 colors will be grouped into 10 clusters. This application is critical also as a pre-processing step for data compression purposes, or to prepare an image data set for other machine learning algorithms.

Then, a new matrix that contains the estimated mean vectors of the clusters assigned to each pixel of the figure is retrieved. This allows us to represent the original picture and the cluster version.

The Python code that implements the image segmentation is also reported in the Appendix and slightly differs from the original one, since some modifications are required in order to deal with an image data input.

The following figures qualitatively show the improvement of the procedure with an increasing number of iterations.




Fig. 4: Original and segmented image, after 10 iterations.

The same convergence diagnostic tools can be implemented in this application as well. In Figure 7, some trace plots of mean vectors are illustrated¹.

¹ Note that the trace plots show value between 0 and 1, rather than between 0 and 255. This is because input data were normalized to 1 and converted in floats prior to training. After that, the cluster means are converted back to an integer scale from 0 to 255.




Fig. 5: Original and segmented image, after 100 iterations.




Fig. 6: Original and segmented image, after 1000 iterations.

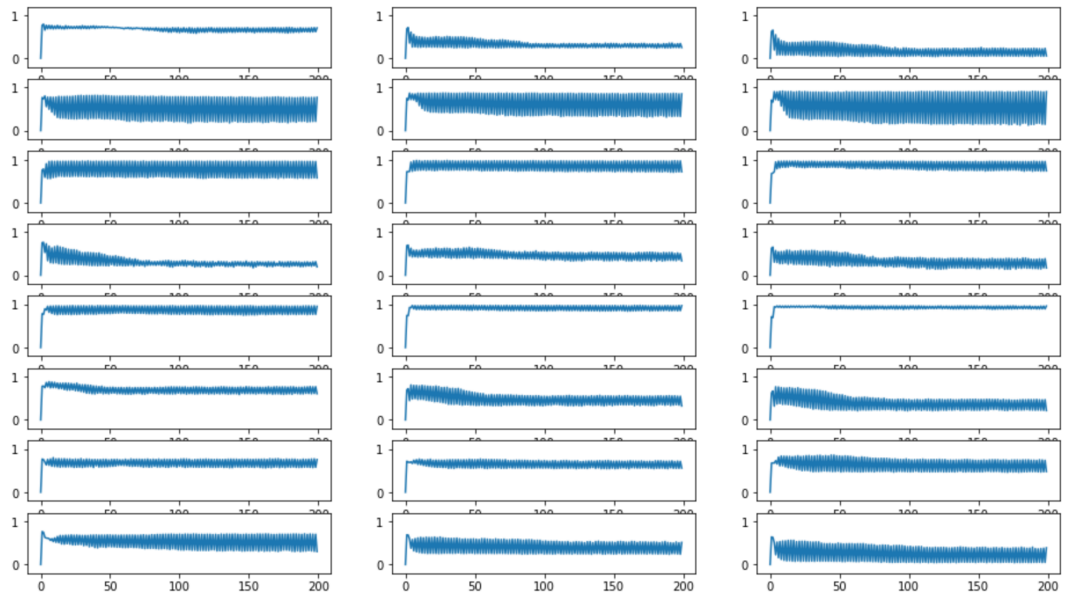


Fig. 7: Trace plots of mean vectors.

6 Python implementation

The algorithm and the applications have been implemented using Python3. The code is also available in the Appendix of this document.

References

1. Bishop, C.M.: Pattern recognition and machine learning. springer (2006)
2. Blei, D.M., Kucukelbir, A., McAuliffe, J.D.: Variational inference: A review for statisticians. Journal of the American statistical Association **112**(518), 859–877 (2017)
3. Casella, G., George, E.I.: Explaining the gibbs sampler. The American Statistician **46**(3), 167–174 (1992)
4. Handcock, M.S., Raftery, A.E., Tantrum, J.M.: Model-based clustering for social networks. Journal of the Royal Statistical Society: Series A (Statistics in Society) **170**(2), 301–354 (2007)
5. Hoff, P.D.: A first course in Bayesian statistical methods, vol. 580. Springer (2009)
6. Johansen, A.M., Evers, L., Whiteley, N.: Monte carlo methods. Lecture notes **200** (2010)
7. Park, J.H.: Color image segmentation using a model-based clustering and a mfa-em algorithm. In: Scandinavian Conference on Image Analysis. pp. 934–941. Springer (2003)
8. Yang, M.H., Ahuja, N.: Gaussian mixture model for human skin color and its applications in image and video databases. In: Storage and retrieval for image and video databases VII. vol. 3656, pp. 458–466. International Society for Optics and Photonics (1998)

Appendix

```

1  # -*- coding: utf-8 -*-
2  #
3  import pandas as pd
4  import numpy as np
5  import random
6  import scipy.stats
7  from scipy.stats import multivariate_normal
8  import matplotlib.pyplot as plt
9
10 #
11 k = 3 # clusters
12 # d = 2 # dimensions
13 prob = 1/k # uniform prior on cluster assignments
14 n = 150 # data points
15 iterations = 1000
16 lambdas = 1
17 sigma = 0.5 # note: sigma^2

```

```

18
19
20 # Generative process (synthetic data)
21 d = 2
22 z_true = np.zeros(shape = (n))
23 z_true = np.random.choice(range(k), p = np.repeat(prob, k), size =
    ↪ n)
24
25 mu_true = np.zeros(shape = (k,d))
26
27 mu_true[0] = np.array((0,0))
28 mu_true[1] = np.array((1,2))
29 mu_true[2] = np.array((3,-4))
30
31 x = np.zeros(shape = (n,d))
32
33 for i in range(n):
34     cluster_assignment = int(z_true[i])
35     x[i] = scipy.stats.multivariate_normal.rvs(mean =
    ↪ mu_true[cluster_assignment], cov = sigma * np.identity(d))
36
37
38
39 # Matrix declaration
40 z = np.zeros(shape = (n, iterations))
41 mu = np.zeros(shape = (k, d, iterations))
42
43 n_k = np.zeros(shape = (k, iterations))
44 k_bar = np.zeros(shape = (k,d, iterations))
45
46
47 # Graphical representation of synthetic data:
48 from pylab import *
49 import matplotlib.pyplot as plt
50
51 x1 = x[:,0]
52 x2 = x[:,1]
53
54 fig = plt.figure(figsize = (16,9))
55 ax = fig.add_subplot(111)
56
57 scatter(x1, x2, s=100 ,marker='o')
58
59 [ plot( [dot_x,dot_x] ,[0,dot_y], ' . ', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x1,x2) ]

```

```

60 [ plot( [0,dot_x] ,[dot_y,dot_y], '.', linewidth = 3 ) for
    ↳ dot_x,dot_y in zip(x1,x2) ]
61
62 left,right = ax.get_xlim()
63 low,high = ax.get_ylim()
64 arrow( left, 0, right -left, 0, length_includes_head = True,
    ↳ head_width = 0.08 )
65 arrow( 0, low, 0, high-low, length_includes_head = True,
    ↳ head_width = 0.08 )
66
67 show()
68
69 # Initialization:
70 mu[:, :, 0] = np.zeros(shape = (k,d)) # mixture locations
71 z[:, 0] = np.random.choice(a = np.array((range(k))), p =
    ↳ np.repeat(prob, k), size = n)
72
73 # Gibbs sampler:
74 for it in range(iterations-1):
75
76     for i in range(n):
77         #cluster_assignment = int(z[i, it])
78         prob_temp = np.zeros(shape = (3,1))
79         for kk in range(k):
80             prob_temp[kk] = prob *
81                 ↳ scipy.stats.multivariate_normal.pdf(x = x[i], mean
82                 ↳ = mu[kk, :, it], cov = sigma * np.identity(d))
83             # print(f"Probabilities: {prob_temp}")
84             prob_temp = prob_temp / prob_temp.sum()
85             # print(f"Probabilities, normalized: {prob_temp}")
86             z[i, it+1] = np.random.choice(a = np.array((0,1,2)), p =
87                 ↳ prob_temp[:, 0])
88
89         for kk in range(k):
90
91             n_k[kk, it] = np.unique(z[:, it],
92                 ↳ return_counts=True)[1][list(np.unique(z[:, 0],
93                 ↳ return_counts=True)[0]).index(kk)]
94
95             for i in range(n):
96                 if z[i, it] == kk:
97                     k_bar[kk, :, it] += x[i]
98             k_bar[kk, :, it] = k_bar[kk, :, it] / n_k[kk, it]
99
100             mean_k = (n_k[kk, it] / sigma) / ((n_k[kk, it]/sigma) +
101                 ↳ 1/lambdas) * k_bar[kk, :, it]

```

```

95         cov_k = 1 / (n_k[kk,it]/sigma + 1/lambdas)
96
97         mu[kk,:,it+1] = scipy.stats.multivariate_normal.rvs(mean =
           ↪ mean_k, cov = cov_k)
98
99         print(f"Iteration {it} complete.")
100
101     # Burn-in and trimming of the relevant chains
102     burn_in = int(iterations/3)
103
104     z_trimmed = z[:,burn_in:]
105     mu_trimmed = mu[:,burn_in:]
106
107     # Bayesian estimates: mean of the approximate posteriors
108     z_est = z_trimmed.mean(axis = 1)
109     mu_est = mu_trimmed.mean(axis = 2)
110
111     # Sampling from the approximate posterior
112     x_sampl_1 = scipy.stats.multivariate_normal.rvs(mean =
           ↪ mu_est[0,:], cov = sigma * np.identity(2), size = 50)
113     x_sampl_2 = scipy.stats.multivariate_normal.rvs(mean =
           ↪ mu_est[1,:], cov = sigma * np.identity(2), size = 50)
114     x_sampl_3 = scipy.stats.multivariate_normal.rvs(mean =
           ↪ mu_est[2,:], cov = sigma * np.identity(2), size = 50)
115
116     # Graphical representation:
117     from pylab import *
118     import matplotlib.pyplot as plt
119
120     x11 = x_sampl_1[:,0]
121     x12 = x_sampl_2[:,0]
122     x13 = x_sampl_3[:,0]
123
124     x21 = x_sampl_1[:,1]
125     x22 = x_sampl_2[:,1]
126     x23 = x_sampl_3[:,1]
127
128     fig = plt.figure(figsize = (16,9))
129     ax = fig.add_subplot(111)
130
131     scatter(x11, x21, s=100 ,marker='o')
132     scatter(x12, x22, s=100 ,marker='o')
133     scatter(x13, x23, s=100 ,marker='o')
134
135     [ plot( [dot_x,dot_x] ,[0,dot_y], ' . ', linewidth = 3 ) for
           ↪ dot_x,dot_y in zip(x11,x21) ]

```

```

136 [ plot( [0,dot_x] ,[dot_y,dot_y], '.', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x11,x21) ]
137
138 [ plot( [dot_x,dot_x] ,[0,dot_y], '.', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x12,x22) ]
139 [ plot( [0,dot_x] ,[dot_y,dot_y], '.', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x12,x22) ]
140
141 [ plot( [dot_x,dot_x] ,[0,dot_y], '.', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x13,x23) ]
142 [ plot( [0,dot_x] ,[dot_y,dot_y], '.', linewidth = 3 ) for
    ↪ dot_x,dot_y in zip(x13,x23) ]
143
144 left,right = ax.get_xlim()
145 low,high = ax.get_ylim()
146 arrow( left, 0, right-left, 0, length_includes_head = True,
    ↪ head_width = 0.08 )
147 arrow( 0, low, 0, high-low, length_includes_head = True,
    ↪ head_width = 0.08 )
148
149 show()
150
151 # Trace plots of the means of three clusters
152 fig, ax = plt.subplots(figsize = (16,9), nrows = k, ncols = d)
153 for ax2 in range(ax.shape[1]):
154     ax[ax1,ax2].plot(mu[0,ax2,:])
155
156 # Effective Sample Size
157 from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
158 from statsmodels.tsa.stattools import acf
159 def eff_sample_size(chain):
160     '''
161     Compute the Effective Sample Size for the MCMC provided
162     '''
163     dimen = len(chain)
164     somma = np.sum(acf(chain))
165     den = 1 + 2*somma
166     ess = dimen/den
167
168     return ess
169
170 # Showing ESS results for cluster means
171 for i in range(k):
172     print(f"ESS (means cluster{i}): ",eff_sample_size(mu[i,0,:]),
    ↪ eff_sample_size(mu[i,0,:]))

```

```

173
174 # Kolmogorov-Smirnov nonparametric test
175 from scipy import stats
176 test_results = np.zeros(50)
177 effective_ss = []
178
179 mu_thinned = mu[:,::2] # thinning
180
181
182 for i in range(k):
183     for j in range(d):
184         burnin = mu_thinned[i,j,:int(len(mu_thinned[i,j,:])/3)]
185         sample1 =
186             ↪ mu_thinned[i,j,int(len(mu_thinned[i,j,:])/3):2*int(len(mu_thinned[i,j,:])/3)]
187         sample2 = mu_thinned[i,j,2*int(len(mu_thinned[i,j,:])/3):]
188         print(f"Test completed for cluster {i}")
189         print(stats.ks_2samp(sample1, sample2))
190
191 z_thinned = z[:,::2]
192 for i in range(n):
193     burnin = z_thinned[i,:int(len(z_thinned)/3)]
194     sample1 =
195         ↪ z_thinned[i,int(len(z_thinned)/3):2*int(len(z_thinned)/3)]
196     sample2 = z_thinned[i,2*int(len(z_thinned)/3):]
197     print(f"Test completed for data point {i}")
198     print(stats.ks_2samp(sample1, sample2))
199
200 # ----- Application 2: image segmentation
201 ↪ -----
202 import pandas as pd
203 import numpy as np
204 import random
205 import scipy.stats
206 from scipy.stats import multivariate_normal
207 import matplotlib.pyplot as plt
208 import math
209
210 #
211 it = 0
212 k = 8 # clusters
213 # d = 2 # dimensions
214 prob = 1/k # uniform prior on cluster assignments
215 # n = 150 # data points

```

```

215 iterations = 200
216 lambdas = 0.003
217 sigma = 0.003 # note: sigma^2
218
219
220 #Image data
221 from matplotlib.image import imread
222 # img = imread('/Users/leonardo/Downloads/xp.jpg')
223 img = imread('home.jpg')
224
225
226 # img = imread('home.png')
227 img_size = img.shape
228 x = img.reshape(img_size[0] * img_size[1], img_size[2])
229 x = x.astype(float)
230
231 d = x.shape[1]
232 n = x.shape[0]
233
234 # Normalizing color intensities between 0 and 1
235 for dim in range(x.shape[1]):
236     for i in range(x.shape[0]):
237         x[i, dim] /= 255
238
239
240 #Matrix declaration
241 z = np.zeros(shape = (n, iterations))
242 mu = np.zeros(shape = (k, d, iterations))
243
244 n_k = np.zeros(shape = (k, iterations))
245 k_bar = np.zeros(shape = (k,d, iterations))
246
247
248 #Gibbs sampler:
249 # mu[:, :, 0] = x.mean(axis = 0) # mixture locations
250 mu[:, :, 0] = np.zeros(shape = (k,d)) # mixture locations
251 z[:, 0] = np.random.choice(a = np.array(range(k)), p =
    ↪ np.repeat(prob, k), size = n)
252
253 for it in range(iterations-1):
254
255     for i in range(n):
256         # cluster_assignment = int(z[i, it])
257         prob_temp = np.zeros(shape = (k,1))
258         for kk in range(k):

```

```

259         prob_temp[kk] = prob *
        ↪ scipy.stats.multivariate_normal.pdf(x = x[i], mean
        ↪ = mu[kk,:,it], cov = sigma * np.identity(d))
260     # print(f"Probabilities: {prob_temp}")
261     prob_temp /= prob_temp.sum()
262     # print(f"Probabilities, normalized: {prob_temp}")
263     z[i,it+1] = np.random.choice(a = np.array(range(k)), p =
        ↪ prob_temp[:,0])
264
265     for kk in range(k):
266         if kk in list(np.unique(z[:,it], return_counts=True)[0]):
267             n_k[kk,it] = np.unique(z[:,it],
        ↪ return_counts=True)[1][list(np.unique(z[:,it],
        ↪ return_counts=True)[0]).index(kk)]
268         for i in range(n):
269             if z[i,it] == kk:
270                 k_bar[kk,:,it] += x[i]
271             k_bar[kk,:,it] = k_bar[kk,:,it] / n_k[kk,it]
272
273         mean_k = (n_k[kk,it] / sigma) / ((n_k[kk,it]/sigma) +
        ↪ 1/lambdas) * k_bar[kk,:,it]
274         cov_k = 1 / (n_k[kk,it]/sigma + 1/lambdas)
275
276         mu[kk,:,it+1] = scipy.stats.multivariate_normal.rvs(mean =
        ↪ mean_k, cov = cov_k)
277
278     print(f"Iteration {it} complete.")
279
280     #Burn-in and trimming
281     burn_in = int(iterations/3)
282     mu_trimmed = mu[:, :, burn_in:]
283     mu_est = mu_trimmed[:, :, :].mean(axis=2)
284
285     x_compr = np.zeros(shape = (1024,d))
286     for i in range(x.shape[0]):
287         x_compr[i] = mu[int(z[i,it-1]), :, it] * 255
288     x_compr = x_compr.astype(int)
289
290     x_reshaped = x.reshape(img_size[0], img_size[1], img_size[2])
291     x_compr = x_compr.reshape(img_size[0], img_size[1], img_size[2])
292
293     #Drawing the final results
294     fig, ax = plt.subplots(1, 2, figsize = (12, 8))
295     ax[0].imshow(img)
296     ax[0].set_title('Original Image')

```



```
297 ax[1].imshow(x_compr)
298 ax[1].set_title(f'Segmented Image with {k} colors')
299 for ax in fig.axes:
300     ax.axis('off')
301 plt.tight_layout();
302
303 #Plot of the means of three clusters
304 fig, ax = plt.subplots(figsize = (16,9), nrows = k, ncols = d)
305 for ax1 in range(ax.shape[0]):
306     for ax2 in range(ax.shape[1]):
307         ax[ax1,ax2].plot(mu[ax1,ax2,:])
308         ax[ax1,ax2].set_ylim((-0.2,1.2))
309
310 #Original pixels, clsuter assignment and assigned cluster means
311 show_its = it-1
312 for i in range(int(x.shape[0]/20)):
313     print(x[i], z[i,show_its], mu_est[int(z[i,show_its]),:])
```